

Collect and validate the videos

Import local clips or collect from YouTube; retain a stable video ID and file hash

Participant assessment

Video assessment

Watch and form an impression

The participant watches the video
In direct or paired mode, they select
their own trust rating from 0 to 10

Prepare the video evidence

Transcribe speech with Whisper
Sample timestamped frames with OpenCV
Link evidence to the registered video file

Explain the judgment

In interview or paired mode, an LLM asks
neutral follow-up questions
Each answer is saved before the next question

Choose the assessment condition

Video evidence alone, or a source record
Keep transcript and frames fixed
Document each condition's inputs

Retain distinct participant measures

Direct rating: the participant's own score
Interview inference: an LLM's interpretation
Insufficient evidence remains unscored

Assess the video independently

LLM trust score from 0 to 10 and reasons
Repeat within each model and temperature
Participant responses are excluded

Compare judgments of the same video

Show score gaps, sample sizes, explanations and variation across model runs
In paired mode, also compare each direct rating with its interview inference